

Lecture 11

Final Exam Review Session

- Basics of Probability
- Basics of Linear Algebra
- Basics of Statistical Inference
- Basics of Information Theory
- Basics of Stochastic Processes
- Inequalities in Information Theory

Contents

- Basics of Probability
- Basics of Linear Algebra
- Basics of Statistical Inference
- Basics of Information Theory
- Basics of Stochastic Processes
- Inequalities in Information Theory

Common functions of random variables

$$\text{prof: } E[(X - E[X])^2] = E[X^2 + (E[X])^2 - 2X E[X]] = E[X^2] + E[(E[X])^2] - 2E[X]E[X]$$

Random Variable X	Discrete	$= E[X^2] + (E[X])^2 - 2E[X]E[X]$	Continuous
Cumulative distribution function (cdf)	$F(a) = P\{X \leq a\}$	$-2E[X]E[X] = E[X^2] - (E[X])^2$	$F(a) = \int_{-\infty}^a f(x)dx$
Probability mass function (pmf) or Probability density function (pdf)	$p(x) = P\{X = x\}$		$f(x) = \frac{d}{dx}F(x)$
Expected value $E[X]$	$\sum_{x: p(x)>0} xp(x)$		$\int_{-\infty}^{\infty} xf(x)dx$
Expected value of $g(x)$, $E[g(x)]$	$\sum_{x: p(x)>0} g(x)p(x)$		$\int_{-\infty}^{\infty} g(x)f(x)dx$
Variance of X , $\text{Var}(X)$		$E[(X - E[X])^2] = E[X^2] - (E[X])^2$	
Standard deviation of X , $\text{std}(X)$			$\sqrt{\text{Var}(X)}$

Common discrete random variables

Name	Probability mass function (pmf)	$E[X]$	$\text{var}(X)$
Uniform	$P(x) = \frac{1}{b-a+1}, \quad x = a, a+1, \dots, b$	$\frac{b+a}{2}$	$\frac{(b-a+1)^2 - 1}{12}$
Binomial	$P(x) = \binom{n}{x} p^x (1-p)^{n-x}, \quad x = 0, 1, \dots, n$	np	$np(1-p)$
Poisson	$P(x) = \frac{e^{-\lambda t} (\lambda t)^x}{x!}, \quad x = 0, 1, \dots, 20$	λt	λt
Geometric	$P(x) = (1-p)^{x-1} p, \quad x = 1, 2, \dots$	$\frac{1}{p}$	$\frac{1-p}{p^2}$
Negative Binomial	$P(x) = \binom{x-1}{r-1} p^r (1-p)^{x-r}, \quad x = r, r+1, \dots$	$\frac{r}{p}$	$\frac{r(1-p)}{p^2}$

rate λt , pmf: $\frac{e^{-\lambda t} (\lambda t)^x}{x!}, \quad x = 0, 1, \dots$

Common continuous random variables

$$\frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

Name	Probability density function (pdf)	$E[X]$	$\text{var}(X)$
Uniform	$\frac{1}{b-a}, \quad a \leq x \leq b$	$\frac{a+b}{2}$	$\frac{(b-a)^2}{12}$
Normal	$\frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}, \quad x \in (-\infty, \infty)$	μ	σ^2
Exponential	$\lambda e^{-\lambda x}, \quad x \geq 0$	$1/\lambda$	$1/\lambda^2$
Gamma	$\frac{\lambda e^{-\lambda x} (\lambda x)^{\alpha-1}}{\Gamma(\alpha)}, \quad x \geq 0, \quad \Gamma(a) = \int_0^\infty e^{-y} y^{a-1} dy$	α/λ	α/λ^2
Beta	$\frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} (1-x)^{\beta-1}, \quad 0 < x < 1$	$\frac{\alpha}{\alpha+\beta}$	$\frac{\alpha\beta}{(\alpha+\beta+1)(\alpha+\beta)^2}$

$f(x) = \lambda e^{-\lambda x}$ is well-defined pdf:

$$\int_0^{+\infty} f(x) dx = \int_0^{+\infty} \lambda e^{-\lambda x} dx = \left[-e^{-\lambda x} \right]_0^{+\infty} = 1$$

Sample Question

1. (20 points) Roll a fair four-sided die twice. Let X be the outcome on the first roll, and Y be the sum of the two rolls. Calculate
- (i) $\mu_X = \mathbb{E}[X]$. (4 points)
 - (ii) $\mu_Y = \mathbb{E}[Y]$. (4 points)
 - (iii) $\sigma_X^2 = \text{Var}(X)$. (4 points)
 - (iv) $\sigma_Y^2 = \text{Var}(Y)$. (4 points)
 - (v) $\text{Cov}(X, Y) := \mathbb{E}[XY] - \mathbb{E}[X]\mathbb{E}[Y]$. (4 points)

Contents

- Basics of Probability
- Basics of Linear Algebra
- Basics of Statistical Inference
- Basics of Information Theory
- Basics of Stochastic Processes
- Inequalities in Information Theory

Definition

- Let A be an $n \times n$ matrix.
- A scalar λ is said to be an **eigenvalue** of A if there exists a nonzero vector x such that

$$Ax = \lambda x$$

eigenvalue
↑
eigenvector associated with λ

- The vector x is said to be an **eigenvector belonging to λ** .

(λ, x) is an eigen-pair of A

Sample Question

2. (20 points) Suppose that A admits eigenvalue decomposition with eigenvalues $\lambda_1, \dots, \lambda_n$. Show that
- (i) $\det(A) = \prod_{i=1}^n \lambda_i$. (5 points)
 - (ii) $\text{Trace}(A) = \sum_{i=1}^n \lambda_i$. (5 points)
 - (iii) The eigenvalues of A^k are $\lambda_1^k, \dots, \lambda_n^k$. (5 points)
 - (iv) The eigenvalues of A^\top are $\lambda_1, \dots, \lambda_n$ as well. (5 points)

Contents

- Basics of Probability
- Basics of Linear Algebra
- **Basics of Statistical Inference**
- Basics of Information Theory
- Basics of Stochastic Processes
- Inequalities in Information Theory

Method of Maximum Likelihood

Suppose that X is a random variable with probability distribution $f(x; \theta)$, where θ is a single unknown parameter. Let x_1, x_2, \dots, x_n be the observed values in a random sample of size n . Then the **likelihood function** of the sample is

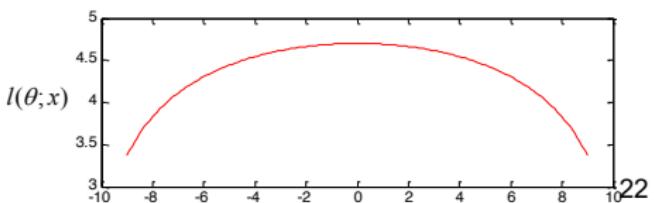
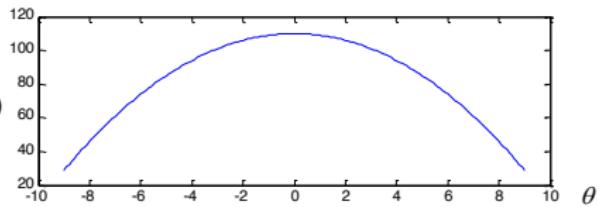
$$L(\theta) = f(x_1; \theta) \cdot f(x_2; \theta) \cdot \dots \cdot f(x_n; \theta) \quad (7-5)$$

Note that the likelihood function is now a function of only the unknown parameter θ . The **maximum likelihood estimator** of θ is the value of θ that maximizes the likelihood function $L(\theta)$.

$$L(\theta; x) = \prod_{i=1}^n f(x_i; \theta) = f(x_1; \theta) \cdots f(x_n; \theta)$$

$$l(\theta; x) = \sum_{i=1}^n \log[f(x_i; \theta)]$$

$$\hat{\Theta}(x) = \arg \max_{\theta} L(\theta; x) = \arg \max_{\theta} l(\theta; x)$$



7-61. A random variable x has probability density function

$$f(x; \theta) = \frac{1}{2\theta^3} x^2 e^{-x/\theta}, \quad 0 < x < \infty, \quad 0 < \theta < \infty$$

Given samples x_1, \dots, x_n ,

find the maximum likelihood estimator for θ

Example: Bernoulli

Let X be a Bernoulli random variable. The probability mass function is

$$f(x; p) = \begin{cases} p^x(1 - p)^{1-x}, & x = 0, 1 \\ 0, & \text{otherwise} \end{cases}$$

where p is the parameter to be estimated. The likelihood function of a random sample of size n is

$$\begin{aligned} L(p) &= p^{x_1}(1 - p)^{1-x_1} p^{x_2}(1 - p)^{1-x_2} \cdots p^{x_n}(1 - p)^{1-x_n} \\ &= \prod_{i=1}^n p^{x_i}(1 - p)^{1-x_i} = p^{\sum_{i=1}^n x_i} (1 - p)^{n - \sum_{i=1}^n x_i} \end{aligned}$$

$$\longrightarrow \ln L(p) = \left(\sum_{i=1}^n x_i \right) \ln p + \left(n - \sum_{i=1}^n x_i \right) \ln(1 - p)$$

$$\longrightarrow \frac{d \ln L(p)}{dp} = \frac{\sum_{i=1}^n x_i}{p} - \frac{\left(n - \sum_{i=1}^n x_i \right)}{1 - p} \longrightarrow \hat{P} = \frac{1}{n} \sum_{i=1}^n X_i$$

Example: normal

Let X be normally distributed with unknown μ and known variance σ^2 . The likelihood function of a random sample of size n , say X_1, X_2, \dots, X_n , is

$$L(\mu) = \prod_{i=1}^n \frac{1}{\sigma\sqrt{2\pi}} e^{-(x_i - \mu)^2/(2\sigma^2)} = \frac{1}{(2\pi\sigma^2)^{n/2}} e^{-(1/2\sigma^2)\sum_{i=1}^n (x_i - \mu)^2}$$

Now

$$\ln L(\mu) = -(n/2) \ln(2\pi\sigma^2) - (2\sigma^2)^{-1} \sum_{i=1}^n (x_i - \mu)^2$$

and

$$\frac{d \ln L(\mu)}{d\mu} = (\sigma^2)^{-1} \sum_{i=1}^n (x_i - \mu)$$

————> What is the MLE for μ ?

Example (Continued, unknown variance)

$$\ln L(\mu, \sigma^2) = -\frac{n}{2} \ln(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2$$

$$\frac{\partial \ln L(\mu, \sigma^2)}{\partial \mu} = \frac{1}{\sigma^2} \sum_{i=1}^n (x_i - \mu) = 0$$

$$\frac{\partial \ln L(\mu, \sigma^2)}{\partial (\sigma^2)} = -\frac{n}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{i=1}^n (x_i - \mu)^2 = 0$$

The solutions to the above equation yield the maximum likelihood estimators

$$\hat{\mu} = \bar{X} \quad \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$$

MLE: Exponential

Let X be a exponential random variable with parameter λ .
The likelihood function of a random sample of size n is:

$$L(\lambda) = \prod_{i=1}^n \lambda e^{-\lambda x_i} = \lambda^n e^{-\lambda \sum_{i=1}^n x_i}$$

$$\ln L(\lambda) = n \ln(\lambda) - \lambda \sum_{i=1}^n x_i$$

$$\frac{d \ln L(\lambda)}{d \lambda} = \frac{n}{\lambda} - \sum_{i=1}^n x_i = 0$$

$$\hat{\lambda} = n \sqrt[n]{\sum_{i=1}^n x_i} = 1/\bar{X} \quad (\text{same as moment estimator})$$

Methods of Moments

Population and samples moments

Let X_1, X_2, \dots, X_n be a random sample from the probability distribution $f(x)$, where $f(x)$ can be a discrete probability mass function or a continuous probability density function. The k th **population moment** (or **distribution moment**) is $E(X^k)$, $k = 1, 2, \dots$. The corresponding k th **sample moment** is $(1/n) \sum_{i=1}^n X_i^k$, $k = 1, 2, \dots$.

Population moments $\mu'_k = \begin{cases} \int x^k f(x) dx & \text{If } x \text{ is continuous} \\ \sum_x x^k f(x) & \text{If } x \text{ is discrete} \end{cases}$

Sample moments $m'_k = \frac{\sum_{i=1}^n X_i^k}{n}$

Method of Moments

- Equating empirical moments to theoretical moments

Let X_1, X_2, \dots, X_n be a random sample from either a probability mass function or probability density function with m unknown parameters $\theta_1, \theta_2, \dots, \theta_m$. The **moment estimators** $\hat{\theta}_1, \hat{\theta}_2, \dots, \hat{\theta}_m$ are found by equating the first m population moments to the first m sample moments and solving the resulting equations for the unknown parameters.

***m* equations for *m* parameters**

$$\begin{cases} m'_1 = \mu'_1 \\ m'_2 = \mu'_2 \\ \vdots \\ m'_m = \mu'_m \end{cases}$$

Sample Question

3. (20 points) A random variable has probability density function

$$f(x; \theta) = \frac{1}{\theta^2} x^{(1-\theta^2)/\theta^2}, \quad 0 < x < 1, 0 < \theta < \infty.$$

Now, given samples X_1, \dots, X_n , derive the maximum likelihood estimator for the parameter θ .

Contents

- Basics of Probability
- Basics of Linear Algebra
- Basics of Statistical Inference
- **Basics of Information Theory**
- Basics of Stochastic Processes
- Inequalities in Information Theory

Understanding Entropy

- Uncertainty in a single random variable

- Can also be written as:

$$H(x) = - \sum_{x \in X} p(x) \log_2 p(x)$$

$$= E_x[-\log_2 p(x)]$$

$$= E_x[\log \frac{1}{p(x)}]$$

$$H(X) = E \left\{ \log \frac{1}{p(X)} \right\}$$

- Intuition: $H = \log(\# \text{of outcomes/states})$

- Entropy is a functional of $p(x)$

- Entropy is a lower bound on the number of bits need to represent a RV. E.g.: a RV that has uniform distribution over 32 outcomes

X	{1, 2, 3, ..., 8}
1	001
2	010
3	011
4	100
5	101
6	110
7	111
8	000

Properties of entropy

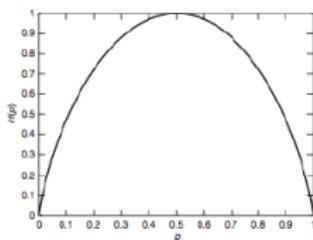
- $H(X) \geq 0$

$$H(X) = E_X[\log \frac{1}{P(X)}] \quad \log \frac{1}{P(X)} \geq 0$$

- Definition, for Bernoulli random variable, $X = 1$ w.p. p ,

$X = 0$ w.p. $1 - p$

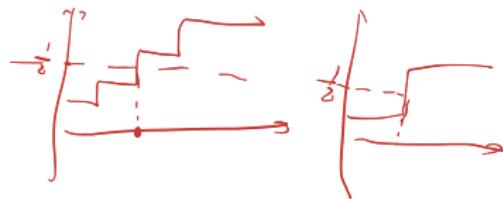
$$H(p) = -p \log p - (1-p) \log(1-p)$$



- **Concave**

- Maximizes at $p = 1/2$

- Example: how to ask questions?



Joint entropy

$$H(X) = -E \log p(x)$$

$$H(Y) = -E \log p(y)$$

$$p(x) \quad p(y)$$

- Extend the notion to a pair of discrete RVs $(X, Y) \Rightarrow H(X, Y) = -E \log p(x, y)$
- Nothing new: can be considered as a single vector-valued RV
- Useful to measure dependence of two random variables

v. $H(X, Y) = - \sum_{x \in \mathcal{X}} \sum_{y \in \mathcal{Y}} p(x, y) \log p(x, y)$

v. $H(X, Y) = -E \log p(X, Y)$

Conditional Entropy

$$(X, Y) \sim p(x, y)$$

$$Y|X \sim p(Y|X)$$

$$Y|X=x \sim p(Y|X=x) = \frac{p(x,y)}{p(x)}, \forall y \in \text{support}(Y)$$

- Conditional entropy: entropy of a RV given another RV. If

$$(X, Y) \sim \underbrace{p(x, y)}$$

$$\textcircled{H(Y|X)} = \sum_{x \in \mathcal{X}} p(x) H(Y|X=x)$$

- Various ways of writing this

$$H(Y|X=x) = - \sum_{y \in \mathcal{Y}} p(y|x) \log p(y|x)$$

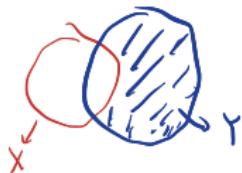
$$H(Y|X) = - \sum_{x \in \mathcal{X}} \sum_{y \in \mathcal{Y}} p(x) p(y|x) \log p(y|x)$$

$$= - \sum_{(x,y) \in \mathcal{X} \times \mathcal{Y}} p(x,y) \log p(y|x),$$

$$= - \mathbb{E} \log p(Y|X)$$

$$H(X, Y)$$

Chain rule for entropy



$$H(Y|X) = H(X, Y) - H(X)$$

- Entropy of a pair of RVs = entropy of one + conditional entropy of the other:

$$H(X, Y) = H(X) + H(Y|X)$$

- Proof:

$$H(X, Y) = - \sum_{(x,y)} p(x,y) \underbrace{\log p(x,y)}$$

$$H(Y|X) \neq H(X|Y)$$

$$= - \sum_{(x,y)} p(x,y) \log p(y|x)$$

$$H(X) - H(X|Y) = H(Y) - H(Y|X)$$

$$\rightarrow H(Y|X)$$

$$= - \sum_{(x,y)} p(x,y) \log p(x) - \sum_{(x,y)} p(x,y) \log p(y|x)$$

$$= - \sum_x \left(\sum_y p(x,y) \right) \log p(x) = - \sum_x p(x) \log p(x) = H(X)$$

Sample Question

4. (20 points) Two random variables X and Y have the following joint distribution:

$$\Pr(X = 0, Y = 0) = 0.2,$$

$$\Pr(X = 0, Y = 1) = 0.3,$$

$$\Pr(X = 1, Y = 0) = 0.1,$$

$$\Pr(X = 1, Y = 1) = 0.4.$$

Calculate $H(X)$, $H(Y)$, $H(X | Y)$, $H(Y | X)$, and $H(X, Y)$.

Contents

- Basics of Probability
- Basics of Linear Algebra
- Basics of Statistical Inference
- Basics of Information Theory
- **Basics of Stochastic Processes**
- Inequalities in Information Theory

Probability and expected time to absorption

- Suppose the state space $\mathcal{S} = \{1, 2, \dots, M\}$
- From state $i \in \mathcal{S}$, denote the probability to reach a specific absorbing state s as a_i .
- It holds that $a_s = 1$ and for all absorbing states $i \neq s$, $a_i = 0$.
- For all transient states i ,

$$a_i = \sum_{j=1}^M a_j p_{ij}.$$



$$a_3 = 1$$

$$a_1 = a_2 P_{12} + a_3 P_{13}$$

$$a_2 = a_1 P_{21} + a_3 P_{23}$$

$$= P_{23} + a_1 P_{21}$$

Probability and expected time to absorption

- Suppose the state space $\mathcal{S} = \{1, 2, \dots, M\}$
- From state $i \in \mathcal{S}$, denote the expected times to absorption as μ_1, \dots, μ_M .
- $\{\mu_i\}_{i \in \mathcal{S}}$ is the unique solution to the system of equations
 - $\mu_i = 0$ for all absorbing state(s) i
 - For all transient states i , $\mu_i = 1 + \sum_{j=1}^M P_{ij} \mu_j$.

$$\mu_i = 1 + \sum_{j=1}^M P_{ij} \mu_j$$

$$\mu_1 = 1 + \sum_{j=1}^M P_{1j} \mu_j$$



$$\mu_3 = 0$$

$$\begin{aligned}\mu_1 &= 1 + \mu_1 P_{11} \\ &\quad + \mu_2 P_{12}\end{aligned}$$

$$\begin{aligned}\mu_2 &= 1 + \mu_1 P_{21} \\ &\quad + \cancel{\mu_3 P_{23}}\end{aligned}$$

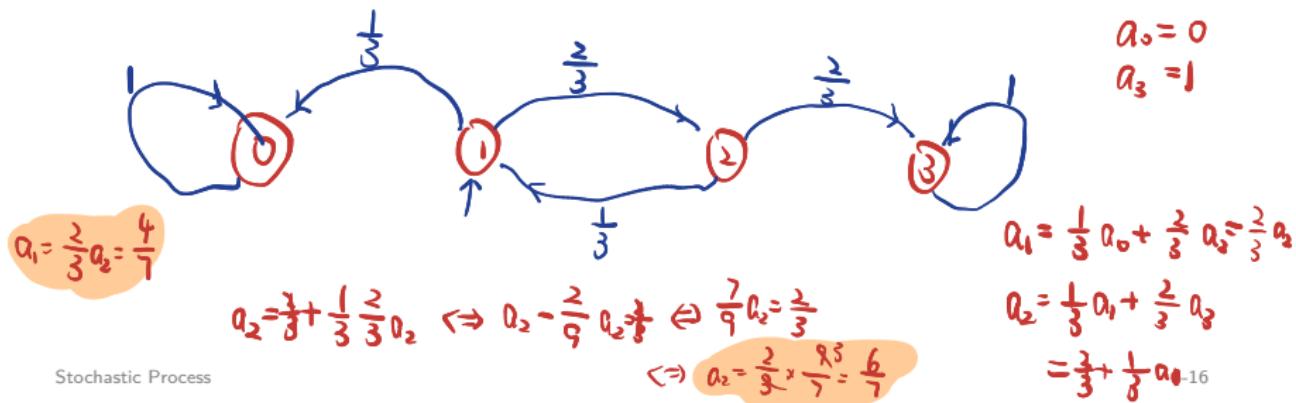
Example

$$S = \{(3,0), (2,1), (1,2), (0,3)\}$$

$$S = \{ \# \text{ of money M have} \} = \{ 0, 1, 2, 3 \}$$

a_i : Prob that M reaches state 3 starting from state i.

Player M has \$1 and player N has \$2. Each game gives the winner \$1 from the other. As a better player, M wins $2/3$ of the games. They play until one of them is bankrupt. What is the probability that M wins?



Sample Question

5. (10 points) Two players, A , and B , start with 2 and 3 dollars respectively. Player A wins each round with probability $p = 0.4$. These two players play such a game until one is ruined.
- (i) Find the probability that A wins all money (ruins B). (5 points)
 - (ii) Compute the expected number of rounds until the game ends. (5 points)

Contents

- Basics of Probability
- Basics of Linear Algebra
- Basics of Statistical Inference
- Basics of Information Theory
- Basics of Stochastic Processes
- Inequalities in Information Theory

Log sum inequality

- Consequence of concavity of \log
- Theorem.** For nonnegative a_1, \dots, a_n and b_1, \dots, b_n

$$\sum_{i=1}^n a_i \log \frac{a_i}{b_i} \geq \left(\sum_{i=1}^n a_i \right) \log \frac{\sum_{i=1}^n a_i}{\sum_{i=1}^n b_i}$$

$$f\left(\frac{a_i}{b_i}\right) = \frac{a_i}{b_i} \log \frac{a_i}{b_i}$$

Equality iff $a_i/b_i = \text{constant}$.

- Proof by Jensen's inequality using convexity of $f(x) = x \log x$.

Application of log-sum inequality

- Very handy in proof: e.g., prove $D(p\|q) \geq 0$:

$$\begin{aligned} D(p\|q) &= \sum_x p(x) \log \frac{p(x)}{q(x)} \\ &\geq \left(\sum_x p(x) \right) \log \frac{\sum_x p(x)}{\sum_x q(x)} = 1 \log 1 = 0. \end{aligned}$$

Convexity of relative entropy

Theorem. $D(p\|q)$ is convex in the pair (p, q) : given two pairs of pdf,

$$D(\lambda p_1 + (1 - \lambda)p_2 \| \lambda q_1 + (1 - \lambda)q_2) \leq \lambda D(p_1 \| q_1) + (1 - \lambda)D(p_2 \| q_2)$$

for all $0 \leq \lambda \leq 1$.

Proof: By definition and log-sum inequality

$$\begin{aligned} & D(\lambda p_1 + (1 - \lambda)p_2 \| \lambda q_1 + (1 - \lambda)q_2) \\ & \underset{x \in \mathcal{X}}{\sum} = (\lambda p_1 + (1 - \lambda)p_2) \log \frac{\lambda p_1 + (1 - \lambda)p_2}{\lambda q_1 + (1 - \lambda)q_2} \Bigg) \text{ log-sum inequality} \\ & \quad \underset{x \in \mathcal{X}}{\sum} \leq \lambda p_1 \log \frac{\lambda p_1}{\lambda q_1} + (1 - \lambda) \log \frac{(1 - \lambda)p_2}{(1 - \lambda)q_2} \\ & \quad = \lambda D(p_1 \| q_1) + (1 - \lambda)D(p_2 \| q_2) \end{aligned}$$

Concavity of entropy

Entropy

$$H(p) = - \sum_i p_i \log p_i$$

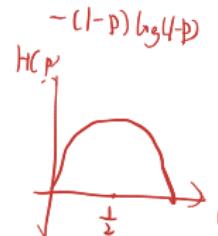
is concave in p

Proof 1:

$$\begin{aligned} H(p) &= - \sum_{i \in \mathcal{X}} p_i \log p_i = - \sum_{i \in \mathcal{X}} p_i \log \frac{p_i}{u_i} u_i \\ &= - \sum_{i \in \mathcal{X}} p_i \log \frac{p_i}{u_i} - \sum_{i \in \mathcal{X}} p_i \log u_i \\ &= -D(p\|u) - \log \frac{1}{|\mathcal{X}|} \sum_{i \in \mathcal{X}} p_i \\ &= \log |\mathcal{X}| - D(p\|u) \end{aligned}$$

$$\mathcal{X} = \{0, 1\}$$

$$\mathbb{P} \quad H(p) = -p \log p$$



Concavity of entropy (Proof 2)

Proof 2: We want to prove

$$H(\lambda p_1 + (1 - \lambda) \underbrace{p_2}) \geq \lambda H(p_1) + (1 - \lambda) H(p_2).$$

A neat idea: introduce auxiliary RV:

$$\theta = \begin{cases} 1, & \text{w.p. } \lambda \\ 2, & \text{w.p. } 1 - \lambda. \end{cases}$$

Let $\underbrace{Z = X_\theta}$, distribution of Z is $\lambda p_1 + (1 - \lambda)p_2$. Conditioning reduces entropy:

$$H(Z) \geq H(Z|\theta)$$

By their definitions

$$H(\lambda p_1 + (1 - \lambda)p_2) \geq \lambda H(p_1) + (1 - \lambda) H(p_2).$$

Concavity and convexity of mutual information

Mutual information $I(X; Y)$ is:

- a concave function of $p(x)$ for fixed $p(y|x)$
- convex function of $p(y|x)$ for fixed $p(x)$

Mixing two gases of equal entropy results in a gas with higher entropy.

Proof of mutual information properties

Proof: write $I(X; Y)$ as a function of $p(x)$ and $p(y|x)$:

$$\begin{aligned} I(X; Y) &= \sum_{x,y} p(x)p(y|x) \log \frac{p(y|x)}{p(y)} \\ &= \sum_{x,y} p(x)p(y|x) \log p(y|x) \\ &\quad - \sum_y \left\{ \sum_x p(x)p(y|x) \right\} \log \left\{ \sum_x p(y|x)p(x) \right\} \end{aligned}$$

- (a): Fixing $p(y|x)$, first linear in $p(x)$, second term concave in $p(x)$
- (b): Fixing $p(x)$, complicated in $p(y|x)$. Instead of verify it directly, we will relate it to something we know.

Strategy for convexity proof

Our strategy is to introduce auxiliary RV \tilde{Y} with a mixing distribution

$$p(\tilde{y}|x) = \lambda p_1(y|x) + (1 - \lambda)p_2(y|x).$$

To prove convexity, we need to prove:

$$I(X; \tilde{Y}) \leq \lambda I_{p_1}(X; Y) + (1 - \lambda)I_{p_2}(X; Y)$$

Since

$$I(X; \tilde{Y}) = D(p(x, \tilde{y}) \| p(x)p(\tilde{y}))$$

We want to use the fact that $D(p \| q)$ is convex in the pair (p, q) .

Completing the convexity proof

What we need is to find out the pdfs:

$$\underbrace{p(\tilde{y})}_{\text{P}(\tilde{y}) = \sum_x \text{P}(y|x)p(x)} = \sum_x [\lambda p_1(y|x)p(x) + (1 - \lambda)p_2(y|x)p(x)] = \lambda p_1(\tilde{y}) + (1 - \lambda)p_2(\tilde{y})$$

We also need

$$p(x, \tilde{y}) = p(\tilde{y}|x)p(x) = \lambda \underbrace{p_1(x, \tilde{y})}_{\text{P}(x, \tilde{y})} + (1 - \lambda) \underbrace{p_2(x, \tilde{y})}_{\text{P}(x, \tilde{y})}$$

Finally, we get, from convexity of $D(p||q)$:

$$\begin{aligned} & \underbrace{D(p(x, \tilde{y}) || p(x)p(\tilde{y}))}_{\text{P}(x, \tilde{y})} \\ &= D(\lambda p_1(y|x)p(x) + (1 - \lambda)p_2(y|x)p(x) || \lambda p(x)p_1(y) + (1 - \lambda)p(x)p_2(y)) \\ &\leq \lambda D(p_1(x, \tilde{y}) || p(x)p_1(y)) + (1 - \lambda) D(p_2(x, \tilde{y}) || p(x)p_2(y)) \\ &= \lambda I_{p_1}(X; Y) + (1 - \lambda) I_{p_2}(X; Y) \end{aligned}$$

Sample Question

6. (10 points) For a discrete probability distribution $P = (p_1, \dots, p_n)$ and parameter $\alpha \in (0, 1)$, the Rényi entropy is defined as

$$H_\alpha(P) = \frac{1}{1-\alpha} \log \left(\sum_{i=1}^n p_i^\alpha \right)$$

- (i) Show that as $\alpha \rightarrow 1$, $H_\alpha(P)$ converges to the Shannon entropy $H(P) = -\sum_{i=1}^n p_i \log p_i$.
(5 points)
- (ii) Show that $H_\alpha(P)$ is concave in p if $\alpha \in (0, 1)$.
(5 points)