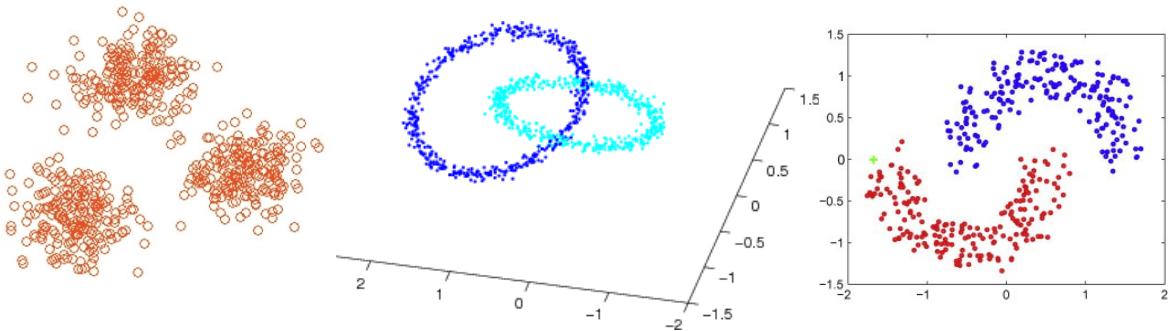


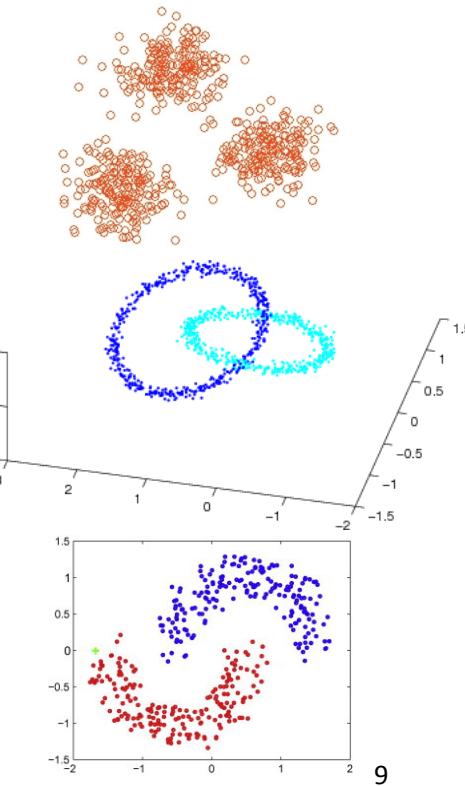
What are the relations
between data points?



Handwritten digits

```
72104149590690159784966540740131  
34727121174235124463556041957843  
74043070291732971627847361369314  
176960549921948713974449254767905  
85665781016467317182029355156034  
46546545144723271818185089250111  
09031642361113952945939036557227  
1284173388792415987230442419577  
28268577918180301994182129759264  
15429204002847124027433003196525  
179316420711215339786361381051315  
56185179462250656372088541140337  
61621928619525442838245031773797  
192119292049148184598837600302661  
95332391268056663882758961841269  
19754089914523789406395213136571  
22632654897130383193446421825488  
40023277687447969098046063548339  
33378087170654380963809968685786  
02402231975108462479309822927359  
18020511376712580371409186774399  
19317397691332336128585114431077  
07944855408215845040615326726931  
46251206217341054311749948402451  
16471942415538314568941538032512  
83440883317359632613607217142821  
79611248177480231310770355276692  
83522560829288887493066321322930  
05781446029147473988471212232323  
91740355868267663279117564951334  
78911691445406223151203812671623  
90152089
```

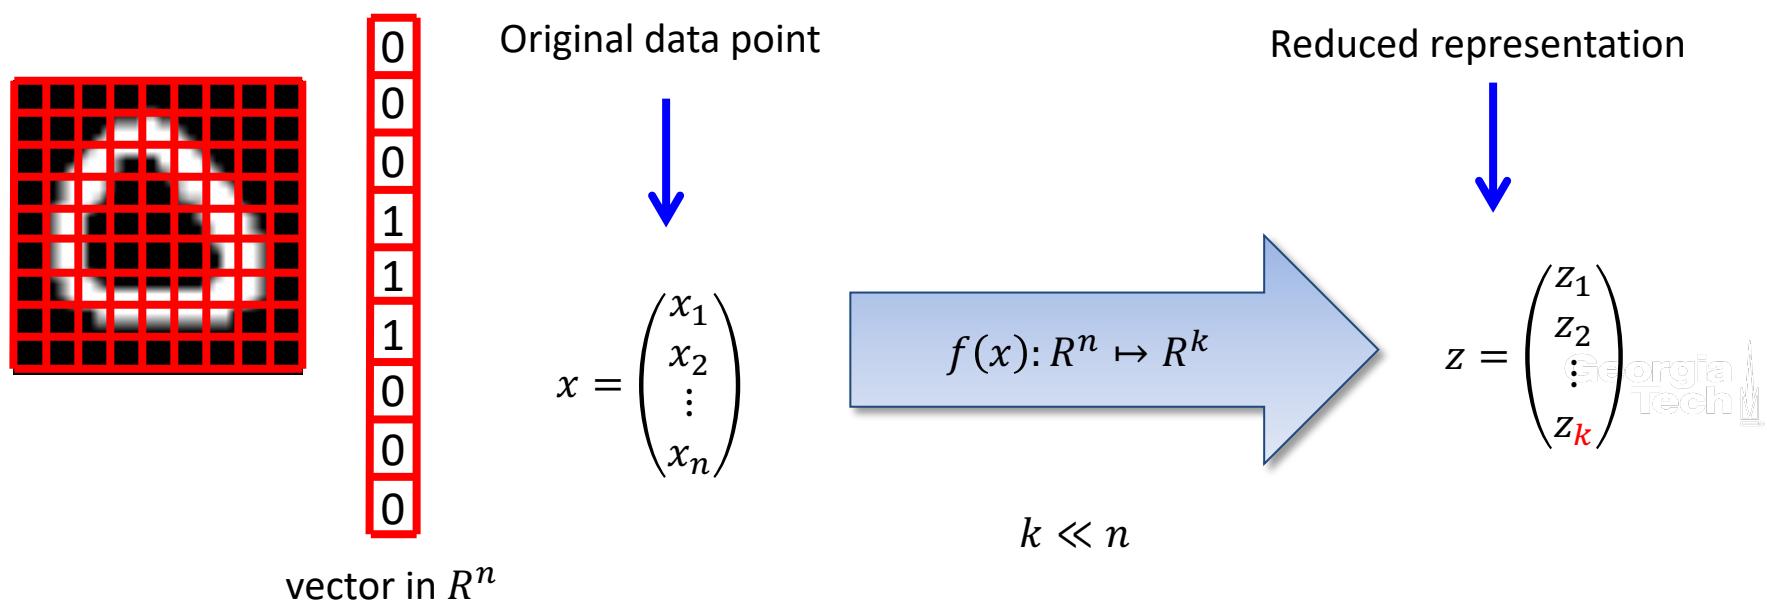
What are the relations between data points?



What is dimensionality reduction?

The process of reducing the number of random variables under consideration

- One can combine, transform or select variables
- One can use linear or nonlinear operations



Why dimensionality reduction and how to think

- The dimension-reduced data can be used for
 - Visualizing, exploring and understanding the data
 - Extracting "features" and dominant modes
 - Cleaning data
 - Speeding up subsequent learning task
 - Building simpler model later
- Applications
 - Image compression
 - Face recognition (eigenface)
 - Natural language processing (latent semantic analysis)

Principal component analysis

Given m data points, $\{x^1, x^2, \dots, x^m\} \in R^d$

Step 1: Estimate the mean and covariance matrix from data

$$\mu = \frac{1}{m} \sum_{i=1}^m x^i \quad \text{and} \quad C = \frac{1}{m} \sum_{i=1}^m (x^i - \mu)(x^i - \mu)^\top$$

Weight vectors

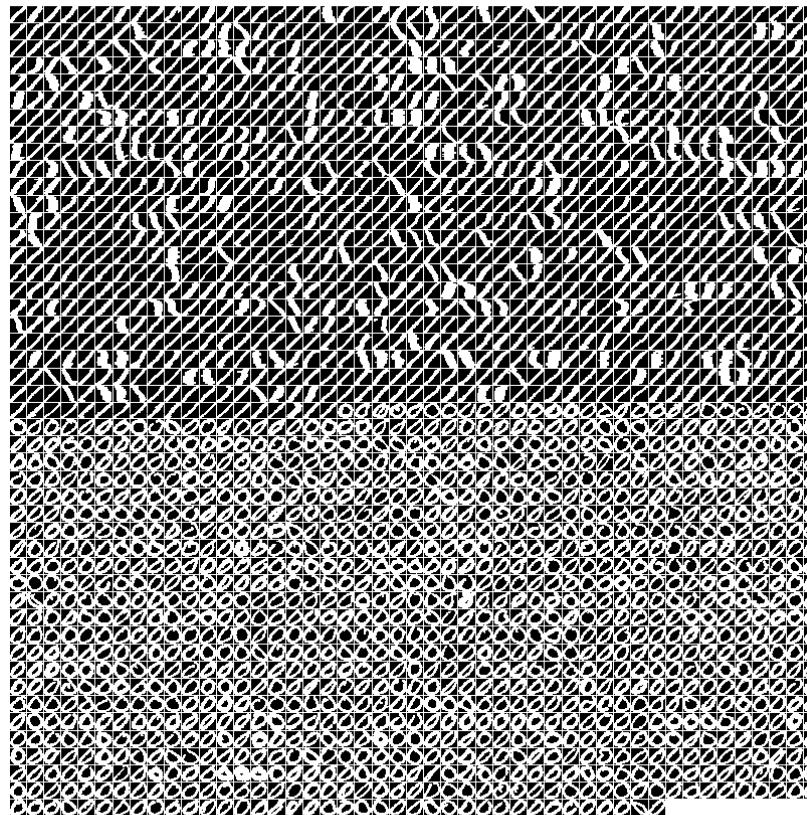
Step 2: Take the eigenvectors w^1, w^2, \dots of C corresponding to the largest eigenvalue λ_1 , the second largest eigenvalue λ_2 ...

Step 3: Compute reduced representation (**principle components** of a data point)

$$z^i = \begin{pmatrix} w^1^\top (x^i - \mu) / \sqrt{\lambda_1} \\ w^2^\top (x^i - \mu) / \sqrt{\lambda_2} \\ \vdots \end{pmatrix}$$

Run demo PCA_digits.m

digit 1 and 0



Run demo PCA_leaf.m



Shape feature	Description
Eccentricity	Eccentricity of the ellipse with identical second moments to I . This value ranges from 0 to 1.
Aspect Ratio	Consider any $X, Y \in \partial I$. Choose X and Y such that $d(X, Y) = D(I)$. Find $Z, W \in \partial I$ maximizing $D^\perp = d(Z, W)$ on the set of all pairs of ∂I that define a segment orthogonal to $[XY]$. The aspect ratio is defined as the quotient $D(I)/D^\perp$. Values close to 0 indicate an elongated shape.
Elongation	Compute the maximum escape distance $d_{\max} = \max_{X \in I} d(X, \partial I)$. Elongation is obtained as $1 - 2d_{\max}/D(I)$ and ranges from 0 to 1. The minimum is achieved for a circular region. Note that the ratio $2d_{\max}/D(I)$ is the quotient between the diameter of the largest inscribed circle and the diameter of the smallest circumscribed circle.
Solidity	The ratio $A(I)/A(H(I))$ is computed, which can be understood as a certain measure of convexity. It measures how well I fits a convex shape.
Stochastic Convexity	This variable extends the usual notion of convexity in topological sense, using sampling to perform the calculation. The aim is to estimate the probability of a random segment $[XY]$, $X, Y \in I$, to be fully contained in I .
Isoperimetric Factor	The ratio $4\pi A(I)/L(\partial I)^2$ is calculated. The maximum value of 1 is reached for a circular region. Curvy intertwined contours yield low values.
Maximal Indentation Depth	Let $C_{H(I)}$ and $L(H(I))$ denote the centroid and arclength of $H(I)$. The distances $d(X, C_{H(I)})$ and $d(Y, C_{H(I)})$ are computed $\forall X \in H(I)$ and $\forall Y \in \partial I$. The indentation function can then be defined as $[d(X, C_{H(I)}) - d(Y, C_{H(I)})]/L(H(I))$, which is sampled at one degree intervals. The maximal indentation depth \mathfrak{D} is the maximum of this function.
Lobedness	The Fourier Transform of the indentation function above is computed after mean removal. The resulting spectrum is normalized by the total energy. Calculate lobedness as $F \times \mathfrak{D}^2$, where F stands for the smallest frequency at which the cumulated energy exceeds 80%. This feature characterizes how lobed a leaf is.

Texture feature	Description
Average Intensity	Average intensity is defined as the mean of the intensity image, m .
Average Contrast	Average contrast is the standard deviation of the intensity image, $\sigma = \sqrt{\mu_2(z)}$.
Smoothness	Smoothness is defined as $R = 1 - 1/(1 + \sigma^2)$ and measures the relative smoothness of the intensities in a given region. For a region of constant intensity, R takes the value 0 and R approaches 1 as regions exhibit larger disparities in intensity values. σ^2 is generally normalized by $(L - 1)^2$ to ensure that $R \in [0, 1]$.
Third moment	μ_3 is a measure of the intensity histogram's skewness. This measure is generally normalized by $(L - 1)^2$ like smoothness.
Uniformity	Defined as $U = \sum_{i=0}^{L-1} p_i^2(z_i)$, uniformity's maximum value is reached when all intensity levels are equal.
Entropy	A measure of intensity randomness.

8 shape features and 6 texture features

Use what criterion for reduction?

There are many criteria (geometric based, information theory based, etc.)

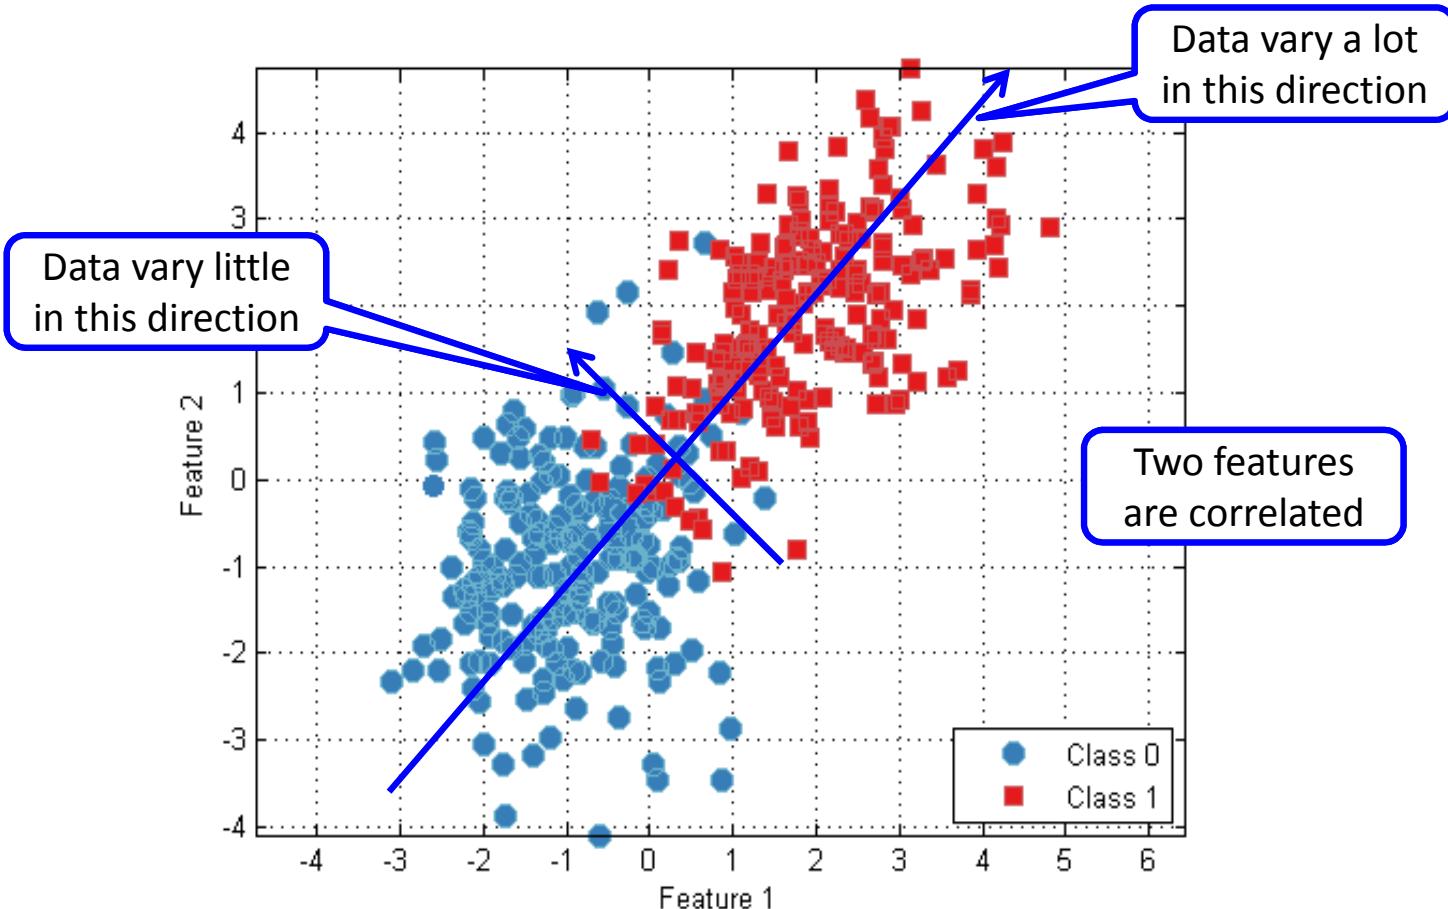
One criterion: want to capture **variation** in data

- variations are “signals” or information in the data
- need to normalize each variables first

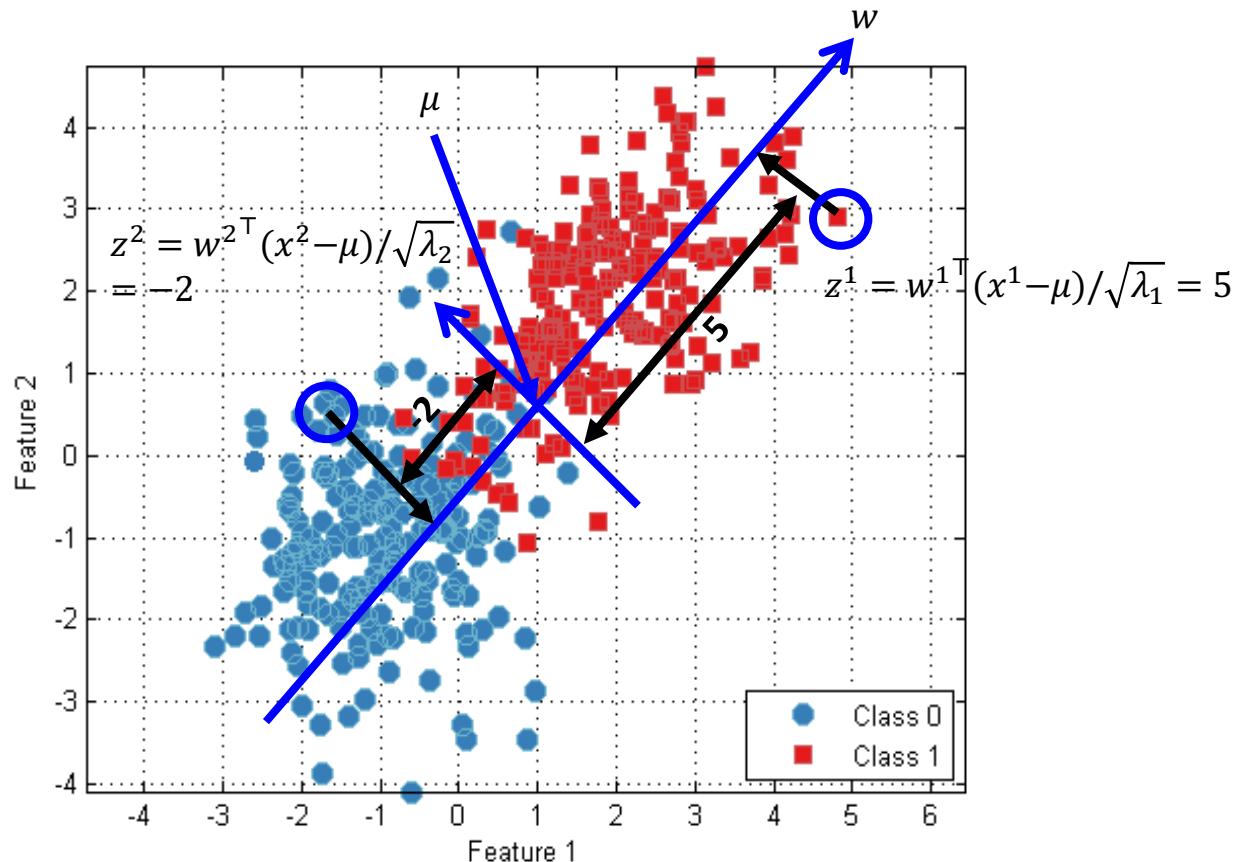
In the process, also discover variables or dimensions highly **correlated**

- represent highly related phenomena
- combine them to form a stronger signal
- lead to simpler presentation

An example



An example (cont.)



How to formulate the problem

Given m data points, $\{x^1, x^2, \dots, x^m\} \in R^n$, with their mean $\mu = \frac{1}{m} \sum_{i=1}^m x^i$

Find a direction $w \in R^n$ where $\|w\| \leq 1$

Such that the variance (or variation) of the data along direction w is maximized

$$\max_{w: \|w\| \leq 1} \frac{1}{m} \sum_{i=1}^m (w^\top x^i - w^\top \mu)^2$$



variance

Is it an easy optimization problem?

Manipulate the objective with linear algebra

$$\begin{aligned} & \frac{1}{m} \sum_{i=1}^m (w^\top x^i - w^\top \mu)^2 \\ &= \frac{1}{m} \sum_{i=1}^m (w^\top (x^i - \mu))^2 \\ &= \frac{1}{m} \sum_{i=1}^m w^\top (x^i - \mu)(x^i - \mu)^\top w \\ &= w^\top \left(\frac{1}{m} \sum_{i=1}^m (x^i - \mu)(x^i - \mu)^\top \right) w \end{aligned}$$

covariance matrix C

Landscape of the optimization problem

Suppose the data has two dimension ($n = 2$)

C is a diagonal matrix

$$C = \begin{pmatrix} 1 & 0 \\ 0 & 2 \end{pmatrix}$$

The optimization problem becomes

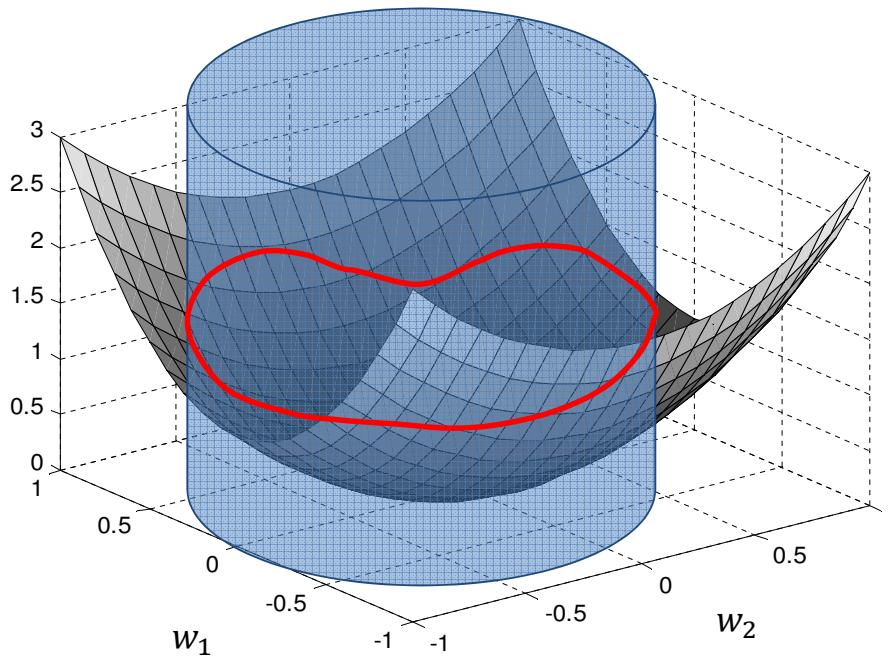
$$\max_{w: \|w\| \leq 1} w^\top C w$$

$$= \max_{w: \|w\| \leq 1} (w_1, w_2) \begin{pmatrix} 1 & 0 \\ 0 & 2 \end{pmatrix} \begin{pmatrix} w_1 \\ w_2 \end{pmatrix}$$

$$= \max_{w: \|w\| \leq 1} w_1^2 + 2w_2^2$$

Landscape of the optimization problem

- $f(w_1, w_2) = w_1^2 + 2w_2^2$



Solving the PCA problem

$$\max_{w: \|w\| \leq 1} w^\top C w, \quad C = \frac{1}{m} \sum_{i=1}^m (x^i - \mu)(x^i - \mu)^\top$$

- Form Lagrangian function of the optimization problem

$$L(w, \lambda) = w^\top C w + \lambda(1 - \|w\|^2)$$

- If w is a maximum of the original optimization problem, then there exists a λ , where (w, λ) is a **stationary point** of $L(w, \lambda)$
- This implies that

$$\frac{\partial L}{\partial w} = 0 = 2Cw - 2\lambda w \Leftrightarrow \textcolor{red}{Cw = \lambda w}$$

- The optimal solution w should be an eigen-vector of C
- Objective function becomes λ (associated with w)

Variance of in the principal direction

- The optimal solution w should be an eigen-vector of C

$$Cw = \lambda w$$

- Objective function becomes λ (associated with w)

$$w^T C w = \lambda w^T w = \lambda \|w\|^2$$

eigen-value

- The problem becomes finding the largest eigenvalue of C

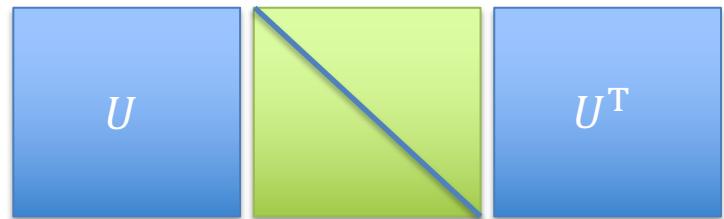
Eigenvalue problem

- Given a symmetric matrix $C \in R^{n \times n}$
 - Find a vector $u \in R^n$ and $\|u\| = 1$
 - Such that

$$Cu = \lambda u$$

- There will be multiple solution: u^1, u^2, \dots, u^n (called the **eigenvectors**) with different $\lambda_1, \lambda_2, \dots, \lambda_n$ (called the **eigenvalues**.)
 - Eigenvectors are ortho-normal: $u^i{}^\top u^i = 1, u^i{}^\top u^j = 0$
 - Eigenvalues are called spectrum
- Eigendecomposition

$$C = U \Lambda U^T$$



Find multiple principal directions

Directions w^1, w^2, \dots which has

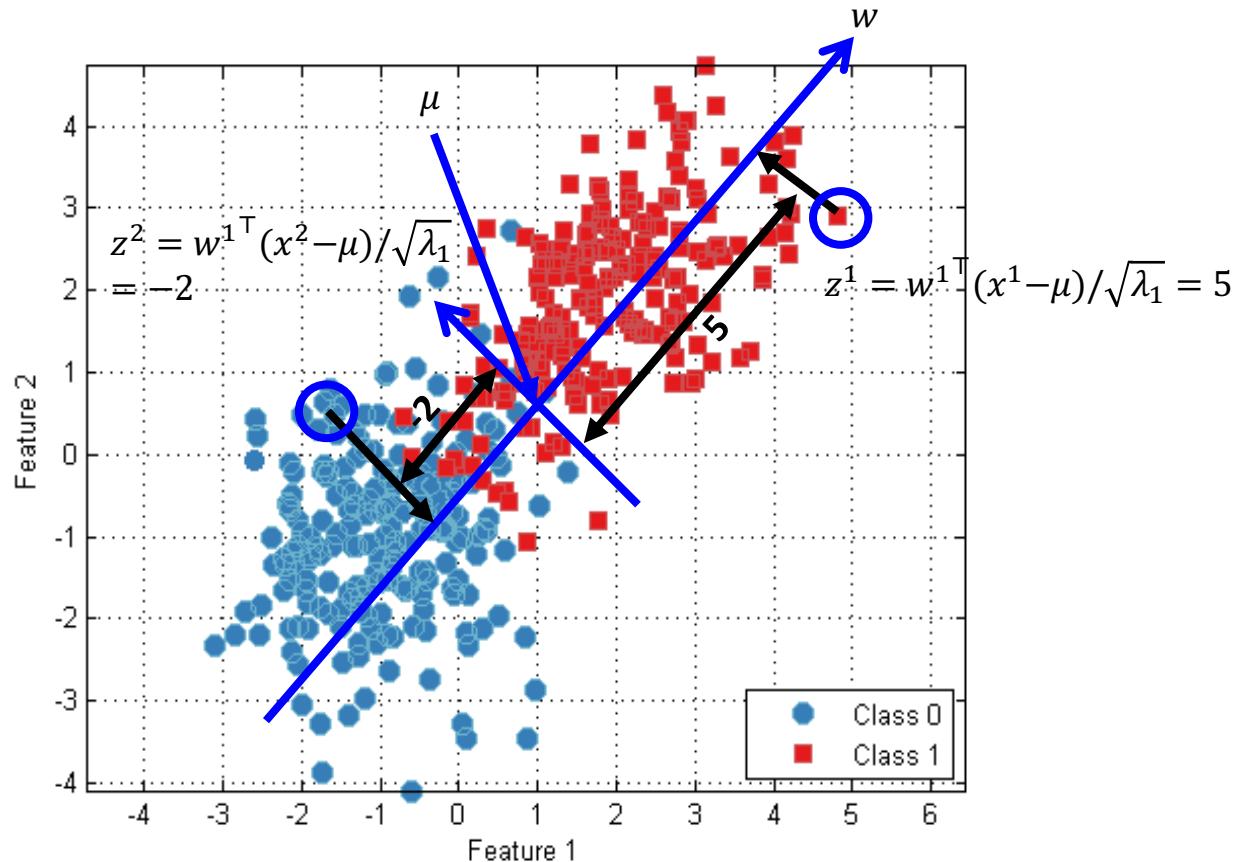
- the largest variances
- **orthogonal** to each other

Take the eigenvectors w^1, w^2, \dots of C corresponding to

- the largest eigenvalue λ_1 ,
- the second largest eigenvalue λ_2

and so on.

An example (cont.)



Principal component analysis

Given m data points, $\{x^1, x^2, \dots, x^m\} \in R^d$

Step 1: Estimate the mean and covariance matrix from data

$$\mu = \frac{1}{m} \sum_{i=1}^m x^i \quad \text{and} \quad C = \frac{1}{m} \sum_{i=1}^m (x^i - \mu)(x^i - \mu)^\top$$

Weight vectors

Step 2: Take the eigenvectors w^1, w^2, \dots of C corresponding to the largest eigenvalue λ_1 , the second largest eigenvalue λ_2 ...

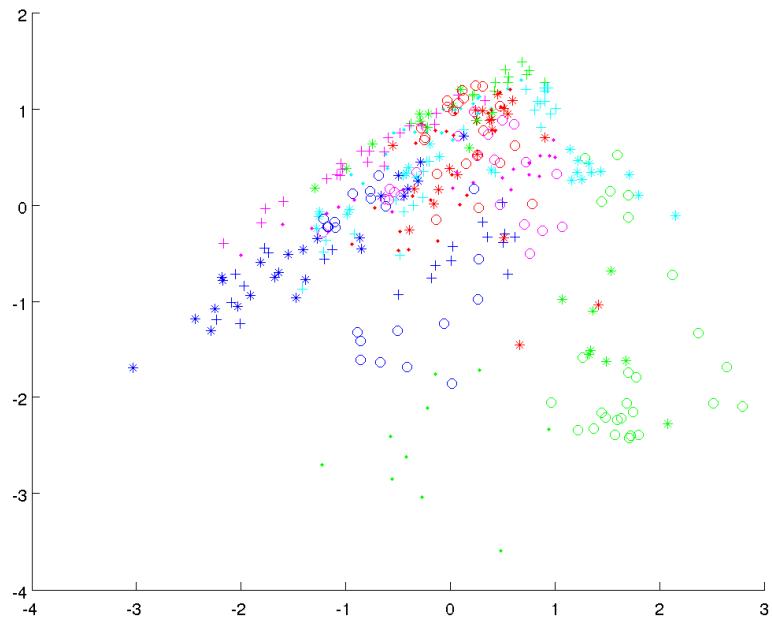
Step 3: Compute reduced representation (**principle components** of a data point)

$$z^i = \begin{pmatrix} w^1^\top (x^i - \mu) / \sqrt{\lambda_1} \\ w^2^\top (x^i - \mu) / \sqrt{\lambda_2} \\ \vdots \end{pmatrix}$$

Look more into PCA_leaf.m



Interpreting the reduced representation



Principal direction:

$W =$

0.0938	0.1924	
0.1902	0.0253	
0.2266	-0.1800	
-0.1850	0.4084	
-0.1600	0.3825	
-0.2063	0.3488	
0.1940	-0.4037	Shape features
0.2150	-0.3566	
-0.3723	-0.2001	
-0.3657	-0.1974	
-0.3602	-0.2037	
-0.3175	-0.1886	
-0.3056	-0.1243	
-0.3482	-0.1829	

Texture
features

Singular Value Decomposition (SVD)

- Singular value decomposition, known as SVD, is a factorization of a real matrix
- For a matrix $M \in \mathbb{R}^{n \times m}$ ($n \leq m$)

$$M = U\Sigma V^\top$$

$$M = [u_1 \ u_2 \ \dots \ u_n] \begin{bmatrix} \sigma_1 & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & \sigma_n \end{bmatrix} [v_1 \ v_2 \ \dots \ v_m]^\top$$

$U \in \mathbb{R}^{n \times n}$
Left singular vectors
(orthonormal)

$\Sigma \in \mathbb{R}^{n \times m}$
Singular values

$V \in \mathbb{R}^{m \times m}$
Right singular vectors
(orthonormal)

where $U \in \mathbb{R}^{n \times n}$, $V^\top \in \mathbb{R}^{m \times m}$, $\Sigma \in \mathbb{R}^{n \times m}$

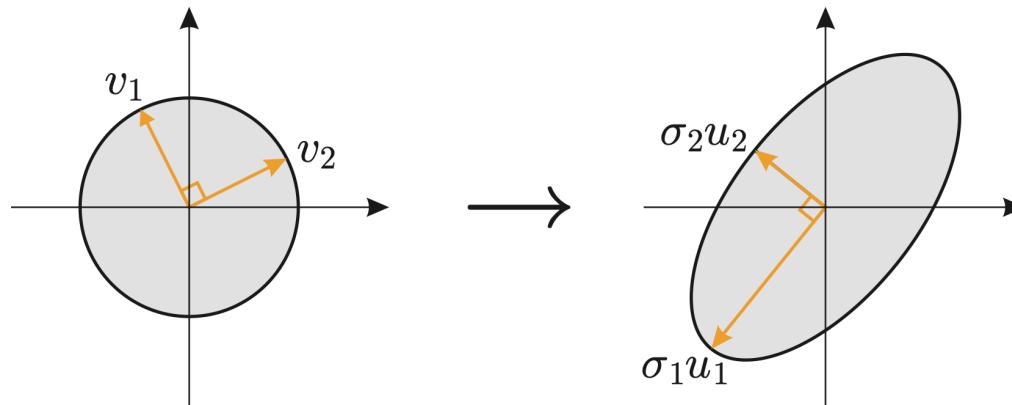
Typically $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_n$

Interpretations of SVD

- A pair of singular vectors (u, v) satisfies

$$Mv = \sigma u \text{ and } M^T u = \sigma v$$

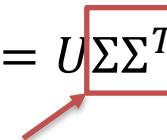
- Geometry



Relationship between SVD and eigendecomposition

$$M = U\Sigma V^T$$

$$C := MM^T = U\Sigma V^T V\Sigma^T U^T = U\Sigma\Sigma^T U^T$$


$$\begin{bmatrix} \sigma_1^2 & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & \sigma_n^2 \end{bmatrix}$$

- The eigenvectors of $C := MM^T$ is U (the left singular vectors of M)
- The eigenvalues of C is σ_i^2 (squared singular values of M)
- Similar results can be derived for $M^T M$

Another way to perform PCA (using SVD)

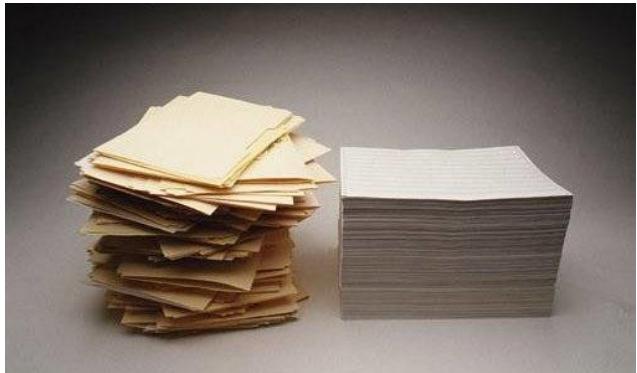
- Note that data covariance matrix can be written as

$$C = \frac{1}{m} \sum_{i=1}^m (x^i - \mu)(x^i - \mu)^T = \frac{1}{m} XX^T$$

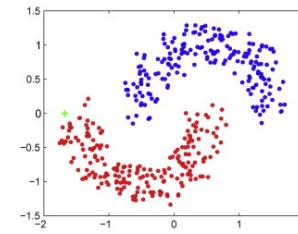
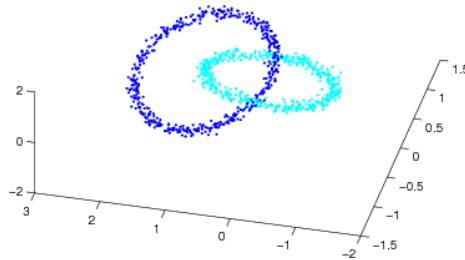
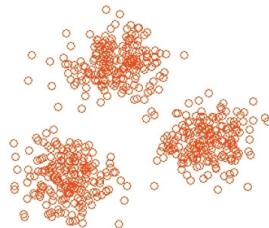
$$X = [x^1 - \mu, \dots, x^m - \mu] \in R^{n \times m}$$

- Eigenvectors of C corresponds to left singular vectors of X
- Find the weight vectors $\{w^1, w^2, \dots, w^r\}$ as the r left singular vectors of the data matrix X (r is the number of principle components)

Documents collections



What are the relations
between data points?



<http://submitpapers.com/misc/docCollection.html>

Bag of words representation

document 1

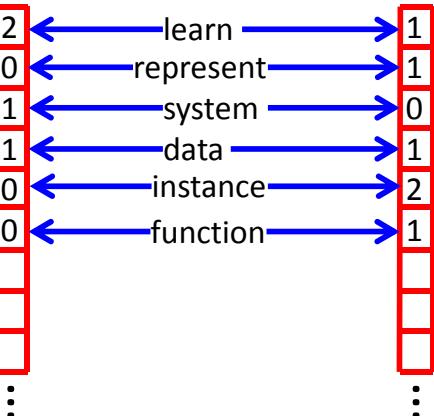
Machine learning concerns the construction and study of systems that can learn from data.

document 2

Representation of data instances and functions evaluated on these instances are part of all machine learning systems

...

vector in R^n



Documents

Feature vector

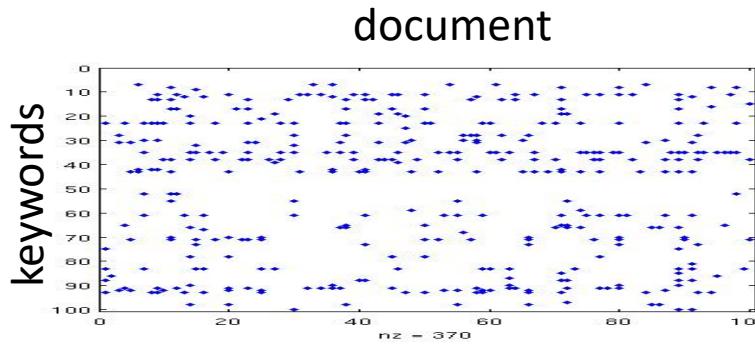
A sample police incident report form from the Atlanta Police Department. The form includes fields for Report Type, Incident Number, Date, Time, Suspect Name, Officer Involved, and various descriptive details about the incident, such as the location (1000 Peachtree Rd NE, Atlanta, GA 30309), time (10:45 AM), and suspect's appearance (Black Male, 30-40 years old). The report is dated 03-22-2007 and is marked as Approved.



Latent semantic analysis (LSA)

- Bag-of-words model or term-document matrix M (more natural language processing techniques: WF/IDF, removing stop words, N-gram)
- Perform PCA of M : (Latent Semantic Analysis, LSA)
- principle components z^i for each document can be interpreted as “feature” of a document

$X =$



Example: Atlanta Police reports

- 20000 police reports, extract 7000 keywords

I Investigator Pickering was advised of a traffic accident near the intersection of Ted Turner Drive and W. Peachtree PI in which the suspects ran from the vehicle. Two individuals were observed by Atlanta Police Officers running from vehicle. One suspect was arrested a short distance from the crash scene.

A free text of a police report

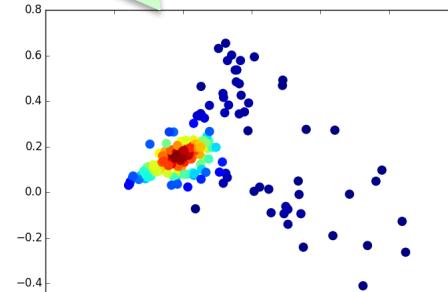
A bag of words for a police report

Terms	Counts
investigator	1
gun	0
kill	0
traffic	1
suspect	2
arrested	1
...	...

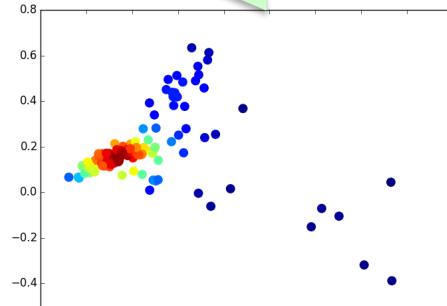
Example: using PCA for data visualization

Atlanta police data, 20000 police reports, 7200 keywords (bi-gram), map into 2 principle components; shown 2d density estimation.

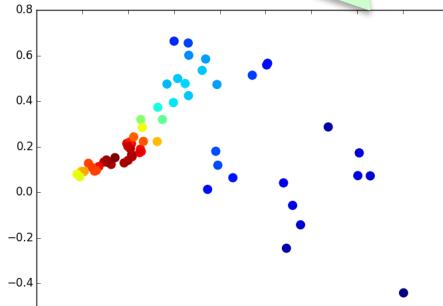
[FRAUD-IMPERS.<\$10,000]



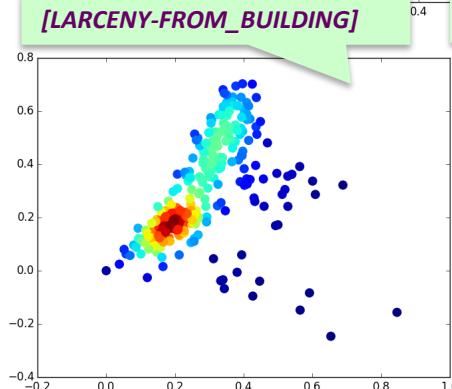
[FRAUD-SWINDLE<\$10,000]



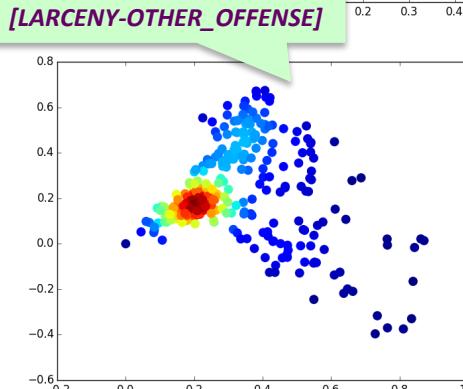
[FRAUD-USE_OF_CRCARD<\$10,000]



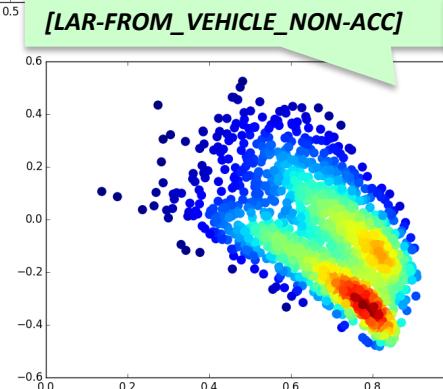
[LARCENY-FROM_BUILDING]



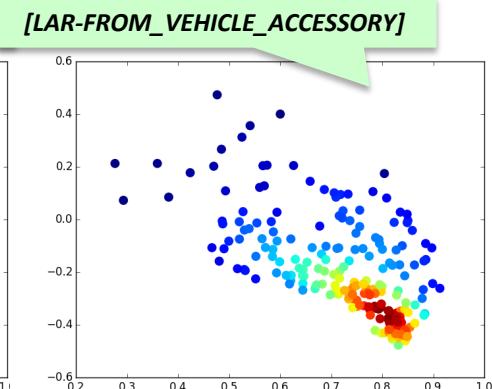
[LARCENY-OTHER_OFFENSE]



[LAR-FROM_VEHICLE_NON-ACC]

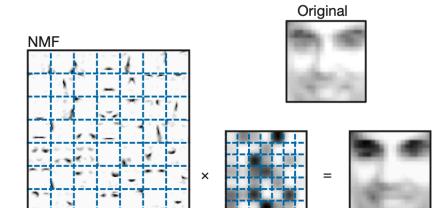


[LAR-FROM_VEHICLE_ACCESSORY]



Extensions/variants of PCA

- Robust PCA: PCA is robust to outliers; it is a common practice to remove outliers to perform PCA; can be cast as a convex optimization problem (Candes, Li, Ma, Wright, 2009)
- Sparse PCA: traditional PCA combines all variables (using the weight vector), not ideal for high-dimensional data; sparse PCA combines just a few important features (solved by optimization) (Johnstone, Yu, 2009)
- Nonlinear PCA: kernelized PCA
- Nonnegative matrix factorization (NMF)
- ICA (independent component analysis): decompose the signal into additive components



Lee, Seung, 1999, Nature.